

# Uniprocessor Scheduling Algorithms

- **Highest Response Ratio:**
  - Aims to minimize *average turn-around time*.
  - Calculate *Normalized Turnaround Time*:
    - $w$  is the total waiting time,  $s$  is the predicted service time.
    - Algorithm for predicting  $s$  value is required.
  - Make priority-queue, prioritize by smallest  $R$  value.
  - Prevents incoming newer processes from starving long, old processes
  - No one starves - since increasing  $w$  value insures  $R$  eventually reaches a sufficiently high value.
- **Multi-level Queue (Feedback):**
  - Two types of task:
    - **Foreground/Interactive:** faster turn-around times are necessary
    - **Background/Batch:** faster turn-around times are preferred
  - Uses multiple **FIFO queues** with different priority, and each has a different time-slice.
    - Higher the priority, fewer the time-slices.
    - Uncompleted higher-priority task pushed to lower-priority queue with more time-slices
  - Process can move down, but they *cannot* move up.
  - Long processes can starve.
- **Guaranteed Scheduling:**
  - Aim is to ensure that a system with  $m$  processes gives  $1/m$  portion of CPU time to each process.
  - $s$  is the actual service time of a process.
  - Run process with *lower score* to ensure  $s/m$  approaches 1.0
- **Lottery:**
  - Each process assigned *Lottery* tickets.
  - Every decision requires choosing a winner randomly.
    - Higher the priority, more the tickets
  - Co-operating processes might *temporarily* exchange tickets
    - Process A request resource from Process B
    - To reduce delay, Process A gives all its tickets to B
    - Process B returns the tickets after receiving resource.
- **Additional Notes:**
  - **Idle Taks:** When nothing scheduled to run, *idle task runs*.
    - Often does '*household*' chores *e.g.* processing statistics, defragmenting disk etc.
  - **Bumping Priority:** High-priority processes often wait for resources being used by low-priority processes.
    - Scheduler can *temporarily* increase priority of the low-priority tasks to minimize high-priority tasks' wait.

# Multiprocessor Scheduling Algorithms

**Multiprocessor Scheduling:** Three types of multiprocessor systems:

- **Distributed:** Relatively autonomous systems that interact.
- **Functionally Specialized:** Specialized devices working on specific domain.
- **Tightly-coupled:** Set of processor sharing main-memory, and being controlled by the OS.
  - **Asymmetric Processing:** Only 'Boss' processor can modify the *Process Control Block*.
    - **Load Balancing:** 'Boss' should aim to uniformly distribute work-load across processors
    - This is at odds with *Processor Affinity*
      - Load balancing *must* consider cost of continuing a *blocked/suspended* process on

non-affinite processor.

- **Symmetric Processing;** All processors can modify *Process Control Block* after using a *Mutex*.

### **Processor Affinity:**

- Processes typically have their own L1, L2, L3, etc. caches.
- Processors possible have *faster* access to certain sections of Main Memory (*NUMA: Non-uniform Memory Access*)
  - Preferable to keep program code in the *same* memory-section
- Desirable to run previously blocked/suspended processes to continue on some processor.

**Multi-core Processor:** Multiple processors mounter atop same chip.

- Minimizes processor-to-processor communication latency

## **Hyper-threading**

**Compute Cycle:** Process is actively utilizing processor.

**Memory Stall:** Process is waiting for a resource (memory access).

**Hyper-threading:** Run two processes on the *same* processor core.

- Run '*Compute Cycle A*' during '*Memory Stall B*'
- Run '*Compute Cycle B*' during '*Memory Stall A*'