

Linux Scheduling (Commercial)

Two types of thread-based schedulers:

- **Linux Real-time Scheduler:** schedules *real-time* and *normal* threads.
 - *normal* threads subject to normal scheduler routine.
- **Linux Non-Real-time Scheduler:** schedules only *normal* threads

Linux Real-time Scheduler

Divides work into 3 queues:

- **SCHED_FIFO:** First-in, first-out real-time threads/tasks.
 - Scheduled as follows:
 - FIFO thread gets interrupted only if:
 - Higher priority FIFO thread is ready.
 - If multiple high priority threads of same priority read, choose one who waited *longer*.
 - FIFO thread gets blocked (e.g on I/O).
 - FIFO yields with `sched_yield`.
 - FIFO thread is interrupted, it is placed in the ready queue.
- **SCHED_RR:** Round-robin Real-time threads/tasks.
 - Same as FIFO scheduling, except implementation includes time-slices.
- **SCHED_OTHER:** Non-real-time threads.
 - A non-real time thread executes *only* if no *Round-robin* or *FIFO* threads are ready.

Linux Non-Real Time Scheduler:

- **Traditional:**
 - $O(n)$ algorithm
 - Bad for *multi-processing* because:
 - Single queue
 - Single *mutex* access to run queue.
 - No *pre-emption*.
- **$O(1)$ Algorithm:**
 - Guarantees $O(1)$ runtime.
 - Implements a queue for 140 different priorities with *round-robin*:
 - 0-99: Real-time; 100-139: Normal
 - A 'queue' is technically two queues:
 - **Active Queue:**
 - Queue of ready, unserved threads.
 - Pre-empted tasks go back to *active queue*
 - *All* scheduling happens here only.
 - When *active queue* empty, *active*, *expired queue* exchange spots.
 - **Expired Queue:** Queue of ready, unserved threads
 - Implements bitmap to indicate empty queues.
- **CFS (Complete Fair Schedule):**
 - Runs in $O(\ln(n))$ time
 - *Caching* used to improve performance.
 - Uses *red-black self-balancing* trees for *ready-queue*.
 - Threads sorted as per *virtual-CPU* running time used
 - **Time Recorded:**
 - *Higher* nice value means over-recorded CPU-runtime

- *Lower* nice value means under-recorded CPU-runtime.
- **History Decay:**
 - *Higher* priority means faster recorded-runtime decay
 - *Lower* priority means slower recorded-runtime decay
- *Round-robin-ing, Time-slice* vary in length to meet *Target Latency*.
 - **Target Latency:** total time to serve all ready processes.
- *Blocked* tasks removed from *ready-queue*
- *Pre-empted* tasks inserted again with updated CPU-runtime
- Trees periodically balanced.
- *I/O Bound* threads highly prioritized to increase *responsiveness*
- **Group Scheduling:**
 - Related threads spawning each other are part of a *group*
 - *Multiple* groups exist on system - groups are treated equally
 - *Groups* might have different number of threads in them
 - Within *groups*, threads have a fair competition

NOTE: you STILL NEED to complete the section about SMP