

I/O Devices and Drivers

For I/O devices, two components are:

- **Bus:** signals for data transfer and communication
- **Controller:** electronics operate on hardware ports, bus, devices etc.

Device Drivers:

- Software that *plugs* into OS interface, tells OS about relevant hardware
- Required to *translate* OS commands to device commands via controller.
- **Windows Support:**
 - Allows drivers to run in *Kernel Mode*
- **UNIX Support:**
 - Able to run *some* drivers in *User Mode*.

Subsystem:

- A *Driver* is connected to the OS via an *interface* called a sub-system.
- *Sub-system* will translate OS commands to *generalized* commands
- *Driver* converts commands to device instructions.

Device Variations:

- Devices can vary in multitude of ways.
 - **Data Transfer Mode:** Devices operate on different data-sizes. *e.g.* keyboard vs. Hard-disk
 - **Access Method:** Devices may provide Sequential, or Random Access
 - **Transfer Schedule:** Devices may communicate, transfer data *Synchronously*, or *Asynchronously*
 - **Dedication:** Devices may be concurrently shareable, or *Dedicated* (*i.e.* no concurrent use).
 - **Device Speed:** Device speeds vary according to device requirements.
- Devices need loose categorization, and then issued only the appropriate commands
 - *e.g.* reading a hard-disk character-by-character is stupid

Data Transfer Mode:

- **Block Interface:**
 - used for devices such as hard-disks, CDs
 - Must support *read*, *write* commands; *seek* commands too if *Random-Access*.
- **Character Interface:**
 - used for devices like keyboards, printers etc. that require *sequential data-transfer*.
 - Useful for devices erratically producing small amounts of data *e.g.* keyboards
 - Must support *get*, *put* commands.

Buffering:

- Area of memory storing data to be transferred.
- Good way to deal with speed mismatch b/w producer and consumer.
- **Double Buffering:**
 - Data arriving while buffer writes to another device needs to be handled
 - Two buffers are used:
 - One for storing incoming data, and another for transferring old data.
 - *e.g.* A keyboard might have two buffers; while one is emptied, second buffer records the keystrokes.
 - Excellent way to handle speed mismatch between producer/consumer

- **Spool:** Buffer-like system where all incoming requests are stored for devices that can only serve one job at a time.

Network Devices:

- Network communication tends to be serial
- Block read commands *read*, *write*, *seek* not appropriate
- *Sockets* model used to communicate b/w client-server:
 - Server opens a *socket*
 - A client can *plugin* to a socket
 - *Packets* of data transferred b/w client and server
- Server uses *select* command to manage a set of clients
 - Displays what *sockets* are waiting to send or receive packets.

I/O Protection:

- All user accesses to I/O *mediated* through OS (since they invoke *system calls*) to confirm validity.
 - e.g. when a user tries to cancel someone's request on the printer
- Extra checking implemented at cost of overhead
- Special circumstances (e.g. high-performance applications and GPUs) might require waiving this condition to improve performance

Kernel I/O Data Structures:

- In Linux, everything is a file
- Kernel has a *global open-file table* i.e. table of any resources under use.
 - Each resource file has pointers to its important functions. e.g. *seek*, *read*, *write*
- Each process records which *files* or *resources* it has 'opened'
 - ...and holds pointer to *files* in the *global open-file table*.

I/O Scheduling:

- Kernel maintains a global structure of I/O requests.
- Kernel re-arranges requests to improve efficiency
- Priority *might* be accounted for; might override the importance of re-arranging it.

I/O Request to Hardware Operations:

Following example illustrates the process:

- Process issues a *read* command
 - Assume resource *file* already open
- Command *invokes* a system call
 - Returns data from *cache* or *buffer* if it's been already read
 - Otherwise, *block* the process, schedule I/O operation, *sub-system* command the *device driver*
- Device-driver uses *buffer* to store incoming data; performs device signalling (i.e. writing into device registers)
- Controller *operates* device as commanded by driver
- Controller will *interrupt* driver or permit *polling* till request is finished.
- *Interrupt Handler* stores the incoming data, notifies *device-driver*, which notifies Kernel
- Kernel transfers data to *address space* of process, process gets un-blocked.
- Scheduler chooses process, execution resumes.