

# Disk Scheduling

- Scheduling algorithms only relevant to Hard-Disks
  - HDDs have non-uniform *random access* times
  - SSDs have constant *random access* time
- **HDD Structure:**
  - **Platter:** Disk composed of *many* platters stacked atop each other.
  - **Track:** Each *platter* divided into circular rings called *track*.
  - **Sector:** Each *track* divided into sections called *sectors*.
  - **Cylinder:** *Tracks* on all platters at the same distance from center form a *Cylinder*.
  - **Arm:** Arm(s) move back-and-forth; used to access *sector* on the corresponding *platter*.
- **Performance Metrics:** HDD performance can be quantified as:
  - **Transfer Rate:** Speed at which data can be *transferred*.
  - **Random Access Time:** Time taken to reach a particular piece of data.
  - **Seek Time:** Time taken to move disk *arm* to the right position.
  - **Rotational Latency:** Time taken to rotate platters to right position
- **NOTE:** *Seek Time* and *Rotation Latency* are generally of about same importance - both are in *milliseconds* range.
- **Average Access Time:** Time taken to *read* data on the disk.
  - For every *track* access, we can calculate the *Average Access Time*.
  - 
  - **T<sub>s</sub>:** Average Seek Time
  - **r:** Rotational Speed
  - **b:** Bytes to be accessed from the *track*
  - **N:** Bytes in the *track*.
- **Sequential Storage:** Data is stored sequentially on same track, on same cylinder.
  - Reading a new *track* in same *cylinder* requires additional *rotational latency*, and *transfer time*
  - Reading a new *tracks* in another *cylinder* requires additional *seek time*, *rotational latency*, and *transfer time*
- **Disk Scheduling Algorithms:**
  - We aim to improve disk performance by optimizing the *Seek Time*.
  - OS is *not* responsible for handling *Rotation Latency*
  - **First come, First serve:**
    - FIFO queue for request storage
    - No optimization attempt made
    - Distance moved by *arm* is high
    - No advantage taken of *spatial locality*
    - No *starvation*
  - **Shortest Seek Time First:**
    - Re-organizes the request queue to begin with closest *track* first
    - *Reduces*, but *doesn't* minimize total distance moved by the *arm*
    - Takes advantage of *spatial locality*.
    - Subject to *starvation* if newly incoming requests closer than older items in request queue
  - **SCAN:**
    - *Arm* moves in one direction, then *reverses* its movement.
      - *Arm* moves till the *end* of the disk reached.
    - Covers less *Arm* distance than *Shortest Seek Time First*
    - No *starvation*
    - Does *not* take advantage of *spatial locality*
  - **C-SCAN (Circular-SCAN):**
    - *Arm* moves in one direction, then *restarts* its movement from the disk's beginning.
    - No *starvation*

- Does *not* take advantage of *spatial locality*
  - Reduces wait for *tracks* at the *periphery*.
    - e.g. assume there are 200 *tracks*, and a periphery-request is on *track* 1.
    - **SCAN:** Assuming that the *scan* happens from the *lower-to-higher count*, and assuming that the original arm position is at *track* 2, disk will take  $2t$  time to reach the periphery-request of low-track number.
    - **C-SCAN:** Time taken will be  $t + s\_max$ , where  $s\_max$  the maximum *seek time*
    - **NOTE:** *seek time* and *scan time* are *NOT* the same.
- **LOOK:**
  - Modification of SCAN and C-SCAN
  - *Arm* changes direction after encountering *last* request in that direction
- **Optimizations:**
  - **Double Bufferring:**
    - SCAN algorithm can be optimized by using *double-bufferring*
    - While the incoming requests are received, previously stored requests are being served
    - Prevents the algorithm-run to be interrupted.
    - Assuming the buffer sizes are  $C$  requests:
      - *Smaller* the buffer, more the algorithm represents the *First-come First Served*.
      - *Larger* the buffer, more the algorithm represents the SCAN algorithm
  - **Real-time Processing:**
    - *Performance* should not be the goal, *predictability* must be.
    - *Higher priority* processes must *not* have to wait for *Lower Priority* if priority difference is too high.
      - Especially important in *Real-time Systems*.